

# A3: Training robust neural networks

Alexandros Kouvatses | Luka Lafaye de Micheaux | Matthieu Neau

# Attack Techniques

- PGD - Projected Gradient Descent
- FGSM - Fast Gradient Sign Method
- MIM - Momentum Iterative Method

# PGD – Projected Gradient Descent

---

## Algorithm 1 PGD Attack

---

```

1: Input: Model  $model$ , images  $X$ , labels  $Y$ 
2: Parameters: Perturbation  $\epsilon$ , step size  $\alpha$ , iterations  $N$ ,
   device  $device$ 
3: Output: Perturbed images  $X'$ 
4:
5:  $X' \leftarrow X$ 
6: Enable gradient computation for  $X'$ 
7: for  $i \leftarrow 1$  to  $N$  do
8:    $outputs \leftarrow model(X')$ 
9:    $loss \leftarrow F.nll\_loss(outputs, Y)$ 
10:  Reset gradients:  $model.zero\_grad()$ 
11:  Compute gradients:  $loss.backward()$ 
12:   $X' \leftarrow X' + \alpha \cdot \text{sign}(X'.grad)$ 
13:  Clip  $X'$  within  $[X - \epsilon, X + \epsilon]$  and  $[0, 1]$ 
14:  Detach  $X'$  from the current graph
15:  Re-enable gradient computation for  $X'$ 
16: end for
17: return  $X'$ 

```

---

## Strengths

- Theoretically grounded in constrained optimization
- Iterative Refinement

## Weaknesses

- Computationally intensive
- Sensitive to hyperparameters
- Overfitting risk

# FGSM

---

**Algorithm 2** FGSM attack

---

```
1: Input: Neural network model, images X, labels Y
2: Parameters: Perturbation  $\epsilon$ , computation device device
3: Output: Perturbed images X'
4:
5: Enable gradient computation for X
6: outputs  $\leftarrow model(X)$ 
7: loss  $\leftarrow F.nll\_loss(outputs, Y)$ 
8: model.zero_grad()
9: Compute gradients: loss.backward()
10: X'  $\leftarrow X + \epsilon \cdot \text{sign}(X.grad)$ 
11: Clip X' to be within valid pixel range [0, 1]
12: return X'
```

---

## Strengths

- Fast with a single step
- Simple to implement

## Weaknesses

- Sensitive to  $\epsilon$
- Mainly designed for  $\ell_\infty$  bounded perturbations

# MIM

---

## Algorithm 3 MIM Attack

---

```

1: Input: Model  $model$ , images  $X$ , labels  $Y$ 
2: Parameters: Max perturbation  $\epsilon$ , step size  $\alpha$ , iterations  $N$ , momentum  $\mu$ , device  $device$ 
3: Output: Adversarially perturbed images  $X'$ 
4:
5: Initialize  $g \leftarrow \mathbf{0}$  (same shape as  $X$ )
6: for  $i \leftarrow 1$  to  $N$  do
7:   Enable gradient computation for  $X'$ 
8:    $outputs \leftarrow model(X')$ 
9:    $loss \leftarrow F.nll\_loss(outputs, Y)$ 
10:  Compute gradients:  $grad \leftarrow$ 
     $autograd.grad(loss, X')[0]$ 
11:  Normalize gradients:  $grad\_norm \leftarrow$ 
     $torch.norm(grad.view(grad.shape[0], -1), p =$ 
     $1, dim = 1)$ 
12:   $grad\_normalized \leftarrow$ 
     $grad / (grad\_norm.view(-1, 1, 1, 1) + 1e - 8)$ 
13:  Update momentum:  $g \leftarrow \mu \cdot g + grad\_normalized$ 
14:   $X' \leftarrow X' + \alpha \cdot sign(g)$ 
15:  Clip change:  $delta \leftarrow torch.clamp(X' - X, min =$ 
     $-\epsilon, max = \epsilon)$ 
16:  Clip  $X'$ :  $X' \leftarrow torch.clamp(X + delta, min =$ 
     $0, max = 1)$ 
17:  Detach  $X'$  from computation graph
18: end for
19: return  $X'$ 

```

---

## Strengths

- Incorporates momentum to stabilize updates
- Reduces oscillations in gradient-based optimization

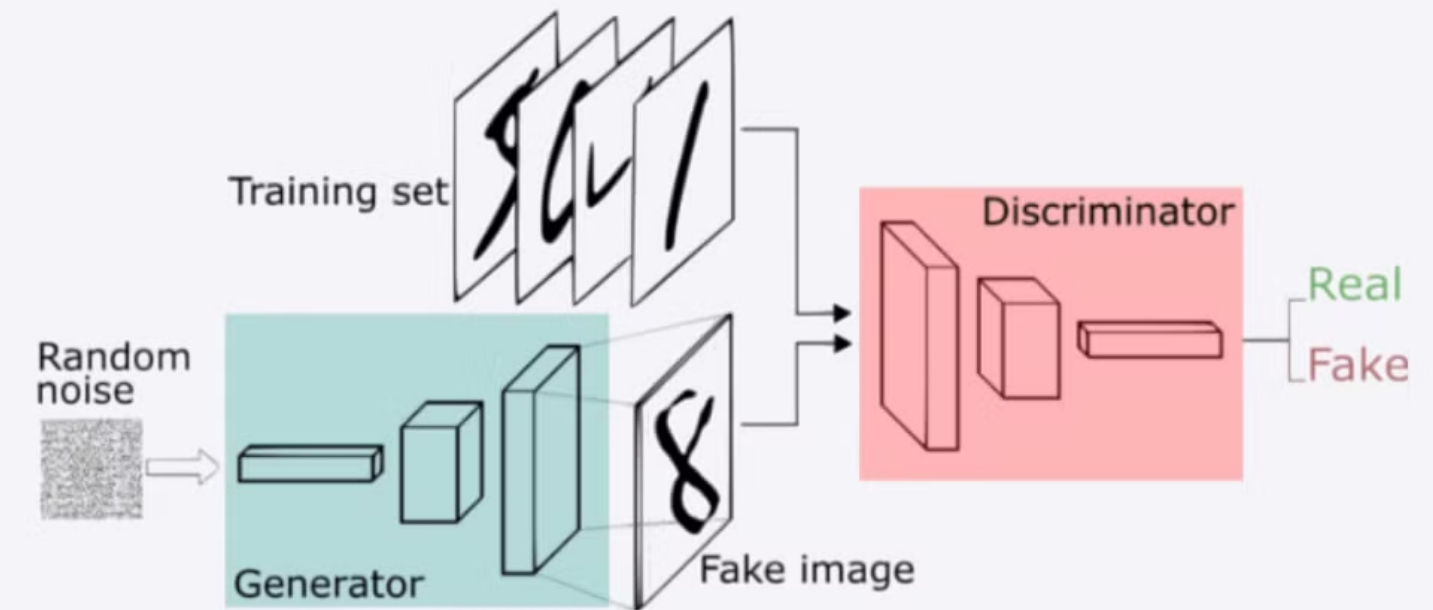
## Weaknesses

- Computationally intensive
- Sensitive to momentum decay



# Defense Techniques

- Adversarial Training - enhances resistance to attacks but reduces clean accuracy
- Regularization Techniques - maintains stable predictions under small perturbations
- Multi-attack technique (adversarial training using all attacks)



$$\mathcal{L}_{\text{total}} = \mathcal{L}(f(x), y) + \lambda \|\nabla_x \mathcal{L}(f(x), y)\|_p^2$$

$\mathcal{L}(f(x), y)$ : Original loss function.

$\|\nabla_x \mathcal{L}(f(x), y)\|_p^2$ : Regularization term penalizing large gradients

# Adversarial Training

- Trains the model on modified data points known as adversarial examples.
- **Objective of Adversarial Examples:** Introduce errors into the model's predictions, aiming to maximize the prediction error during training to make the model more robust.
- **Balanced Training Approach:** Adjusts the training process to include a mix of both natural and adversarially altered inputs, with the aim of minimizing overall prediction errors and enhancing model resilience.

# Regularization Techniques (1)

- **Spectral Normalization:** This technique adjusts each layer's weights by dividing them by their largest singular value, effectively moderating the layer's sensitivity to input perturbations.

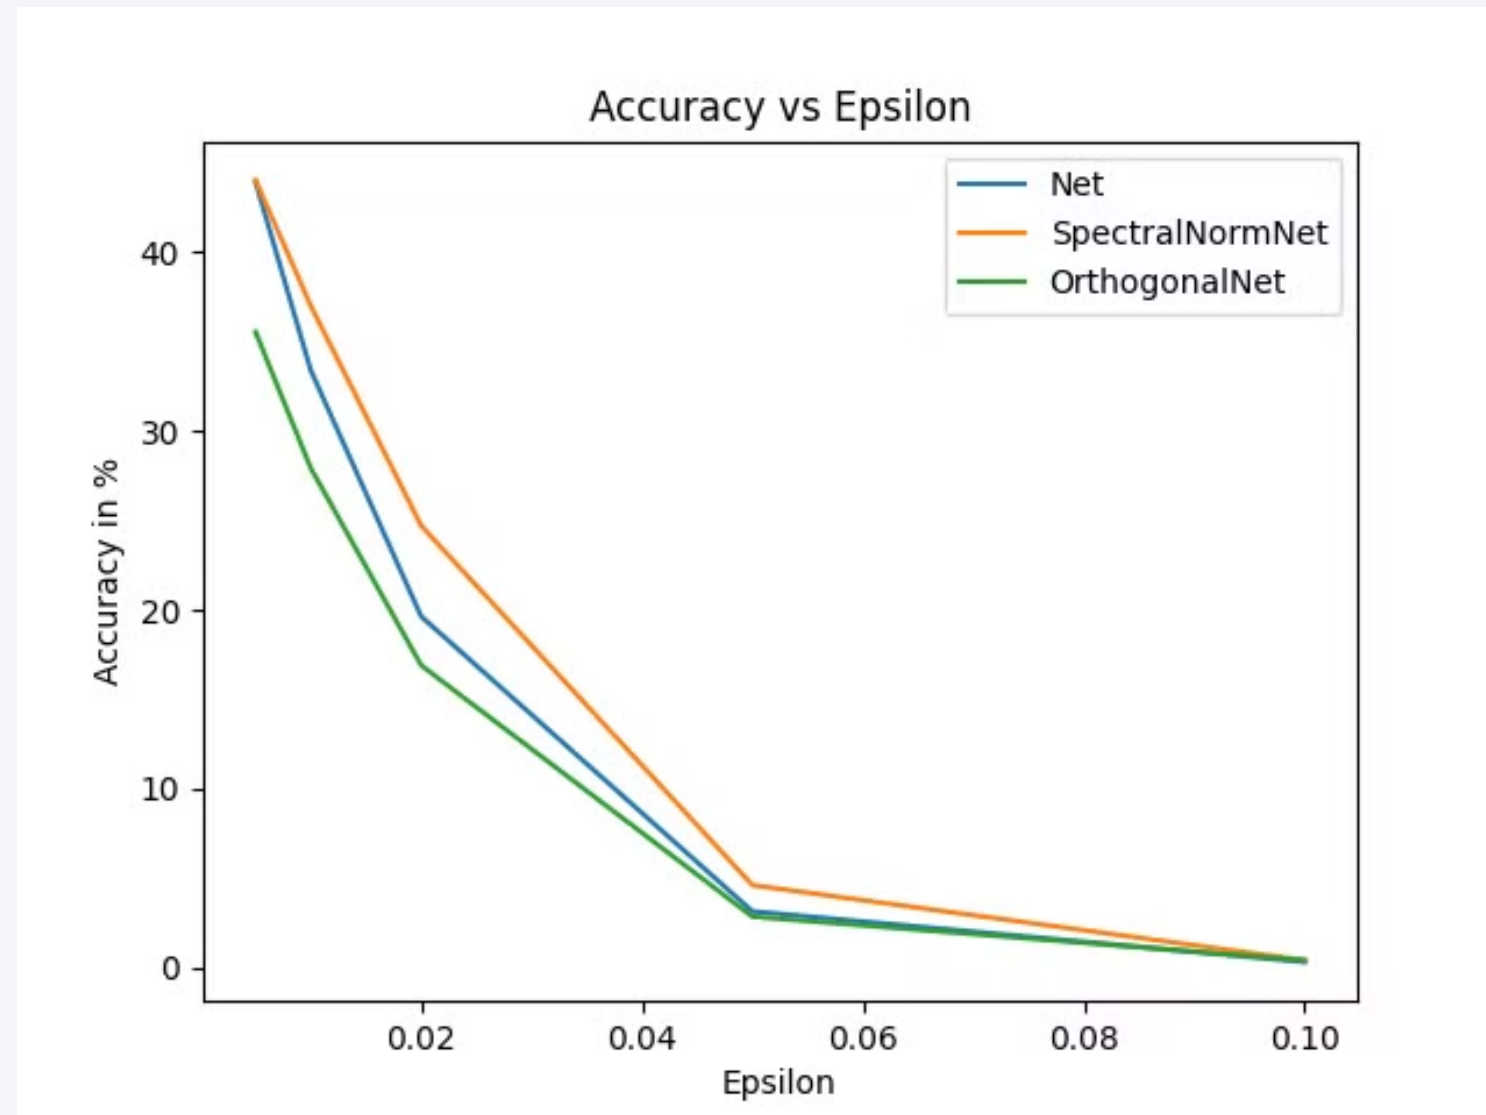
$$\sigma(a) = \max_{\mathbf{h}:\mathbf{h}\neq 0} \frac{\|A\mathbf{h}\|_2}{\|\mathbf{h}\|_2} = \max_{\|\mathbf{h}\|_2\leq 1} \|A\mathbf{h}\|_2 \qquad \bar{W}_{\text{SN}}(W) = W/\sigma(W)$$

- **Orthogonal Normalization:** Attempts to maintain weight matrices close to orthogonality to stabilize learning, though found less effective for tasks unrelated to disentangling latent spaces.

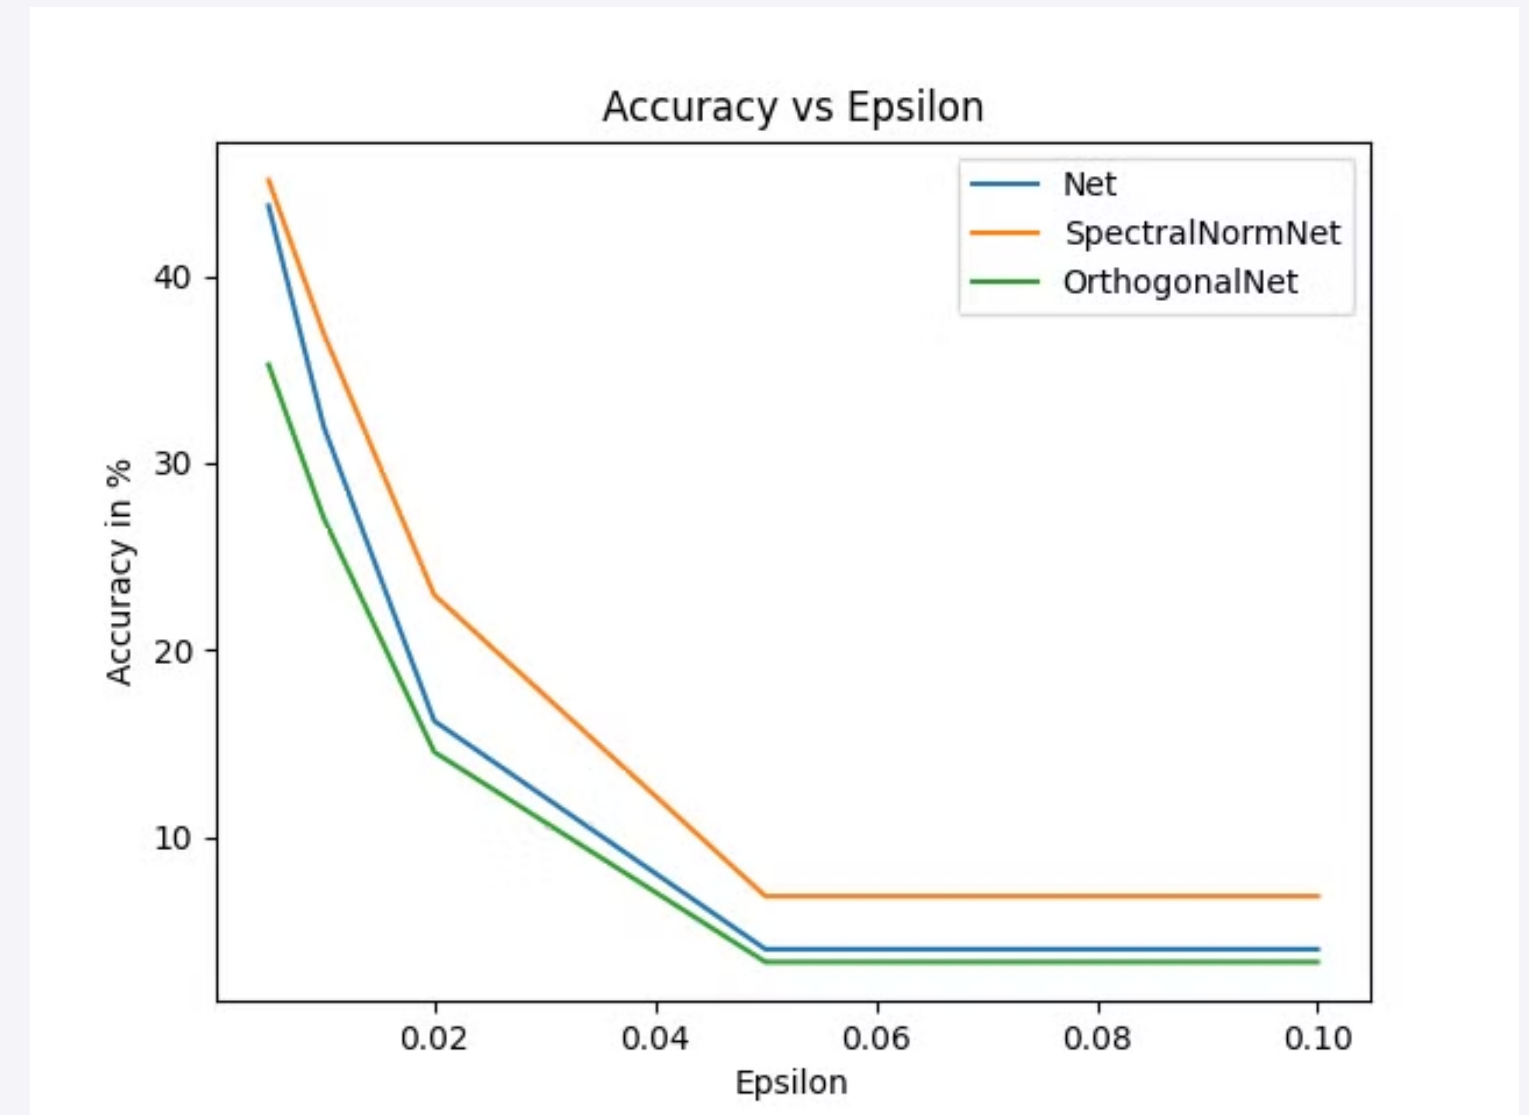
$$\mathcal{L}_{ortho} = \sum (|WW^T - I|)$$



# Regularization Results (1)



FGSM attacks



PGD attacks,  $\alpha = 0.001$ ,  $n_{\text{iter}} = 40$

# Regularization Techniques (2)

- **Gradient Regularization:** Combines adversarial training with a gradient penalty to counteract the effects of input perturbations, improving model robustness against adversarial attacks

$$P = \frac{1}{N} \sum_{i=1}^N \left( \left\| \nabla_{\mathbf{x}'_i} L \right\|_2 \right)^2$$

# Multi-Attack

- **Multi-Attack Technique:** Multiple adversarial attacks (FGSM, MIM, PGD) are applied in a cyclical manner during training, enhancing model robustness by exposing it to varied perturbations.
- **Cyclical Application:** Each training batch applies a different attack based on the batch index, ensuring even exposure to all attack types throughout the training epochs.
- **Diverse Input Utilization:** Utilizing adversarial examples from different attacks as inputs for training prevents model overfitting and contributes to superior performance by increasing input diversity.

# Results

Attack	Defense	Nat Accuracy	PGD l <sub>inf</sub>	PGD l <sub>2</sub>
PGD	Adversarial Training	56.25%	17.26%	26.4%
FGSM	Adversarial Training	50%	0.46%	8.48%
MIM	Adversarial Training	37.5%	17.23%	26.36%
PGD	Gradient Regularization	25%	27.43%	27.01%
PGD/MIM/FGSM	Multi-Attack	43.75%	28.39%	34.87%

**Thank you for  
listening!**

